

```
In [1]: # Importing libraries
import pandas as pd
from IPython.display import Markdown, display
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns; sns.set(style="ticks", color_codes=True)

def printmd(string):
    display(Markdown(string))
```

```
In [2]: data = pd.read_csv('data.csv')
```

# 1. Preliminary Data Analysis

```
In [3]: # Setting all the categorical columns to type category
for col in set(data.columns) - set(data.describe().columns):
    data[col] = data[col].astype('category')

printmd('## 1.1. Columns and their types')
print(data.info())
```

## 1.1. Columns and their types

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 32 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                     569 non-null    int64
1   diagnosis                             569 non-null    category
2   radius_mean                           569 non-null    float64
3   texture_mean                           569 non-null    float64
4   perimeter_mean                         569 non-null    float64
5   area_mean                             569 non-null    float64
6   smoothness_mean                       569 non-null    float64
7   compactness_mean                      569 non-null    float64
8   concavity_mean                        569 non-null    float64
9   concave points_mean                   569 non-null    float64
10  symmetry_mean                         569 non-null    float64
11  fractal_dimension_mean                569 non-null    float64
12  radius_se                             569 non-null    float64
13  texture_se                             569 non-null    float64
14  perimeter_se                           569 non-null    float64
15  area_se                               569 non-null    float64
16  smoothness_se                         569 non-null    float64
17  compactness_se                        569 non-null    float64
18  concavity_se                          569 non-null    float64
19  concave points_se                     569 non-null    float64
20  symmetry_se                           569 non-null    float64
21  fractal_dimension_se                  569 non-null    float64
22  radius_worst                          569 non-null    float64
23  texture_worst                         569 non-null    float64
24  perimeter_worst                       569 non-null    float64
25  area_worst                            569 non-null    float64
26  smoothness_worst                      569 non-null    float64
27  compactness_worst                     569 non-null    float64
28  concavity_worst                       569 non-null    float64
29  concave points_worst                  569 non-null    float64
30  symmetry_worst                        569 non-null    float64
31  fractal_dimension_worst                569 non-null    float64
dtypes: category(1), float64(30), int64(1)
memory usage: 138.6 KB
None
```

```
In [4]: # Top 5 records
printmd('## 1.2. Data')
data.head()
```

## 1.2. Data

Out[4]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_me
0	842302	M	17.99	10.38	122.80	1001.0	0.118
1	842517	M	20.57	17.77	132.90	1326.0	0.084
2	84300903	M	19.69	21.25	130.00	1203.0	0.109
3	84348301	M	11.42	20.38	77.58	386.1	0.142
4	84358402	M	20.29	14.34	135.10	1297.0	0.100

5 rows × 32 columns

```
In [5]: printmd('## 1.3. Summary Statistics')


data.describe()
```

### 1.3. Summary Statistics

Out[5]:

	id	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
count	5.690000e+02	569.000000	569.000000	569.000000	569.000000	569.000000
mean	3.037183e+07	14.127292	19.289649	91.969033	654.889104	0.096360
std	1.250206e+08	3.524049	4.301036	24.298981	351.914129	0.014060
min	8.670000e+03	6.981000	9.710000	43.790000	143.500000	0.052630
25%	8.692180e+05	11.700000	16.170000	75.170000	420.300000	0.086370
50%	9.060240e+05	13.370000	18.840000	86.240000	551.100000	0.095870
75%	8.813129e+06	15.780000	21.800000	104.100000	782.700000	0.105300
max	9.113205e+08	28.110000	39.280000	188.500000	2501.000000	0.163400

8 rows × 31 columns



```
In [6]: printmd('## 1.4. Missing values')

data.isnull().sum()
```

### 1.4. Missing values

```
Out[6]: id                                0
        diagnosis                        0
        radius_mean                      0
        texture_mean                     0
        perimeter_mean                   0
        area_mean                        0
        smoothness_mean                  0
        compactness_mean                 0
        concavity_mean                   0
        concave points_mean              0
        symmetry_mean                    0
        fractal_dimension_mean           0
        radius_se                        0
        texture_se                       0
        perimeter_se                     0
        area_se                          0
        smoothness_se                   0
        compactness_se                   0
        concavity_se                     0
        concave points_se                0
        symmetry_se                      0
        fractal_dimension_se             0
        radius_worst                     0
        texture_worst                    0
        perimeter_worst                  0
        area_worst                       0
        smoothness_worst                 0
        compactness_worst                0
        concavity_worst                  0
        concave points_worst             0
        symmetry_worst                   0
        fractal_dimension_worst          0
dtype: int64
```

```
In [7]: printmd('## 1.5. Clean Data')

# Cleaning and modifying the data - dropping 'id' cloumn/attributues
data.drop('id',axis=1, inplace=True)
data.head()
```

## 1.5. Clean Data

Out[7]:

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compa
0	M	17.99	10.38	122.80	1001.0	0.11840	
1	M	20.57	17.77	132.90	1326.0	0.08474	
2	M	19.69	21.25	130.00	1203.0	0.10960	
3	M	11.42	20.38	77.58	386.1	0.14250	
4	M	20.29	14.34	135.10	1297.0	0.10030	

5 rows × 31 columns

```
In [8]: printmd('## 1.6. Correlation Matrix')

display(data.corr())

printmd('We see that none of the columns are highly correlated.')
```

## 1.6. Correlation Matrix

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_me
radius_mean	1.000000	0.323782	0.997855	0.987357	0.1705
texture_mean	0.323782	1.000000	0.329533	0.321086	-0.0233
perimeter_mean	0.997855	0.329533	1.000000	0.986507	0.2072
area_mean	0.987357	0.321086	0.986507	1.000000	0.1770
smoothness_mean	0.170581	-0.023389	0.207278	0.177028	1.0000
compactness_mean	0.506124	0.236702	0.556936	0.498502	0.6591
concavity_mean	0.676764	0.302418	0.716136	0.685983	0.5219
concave points_mean	0.822529	0.293464	0.850977	0.823269	0.5536
symmetry_mean	0.147741	0.071401	0.183027	0.151293	0.5577
fractal_dimension_mean	-0.311631	-0.076437	-0.261477	-0.283110	0.5847
radius_se	0.679090	0.275869	0.691765	0.732562	0.3014
texture_se	-0.097317	0.386358	-0.086761	-0.066280	0.0684
perimeter_se	0.674172	0.281673	0.693135	0.726628	0.2960
area_se	0.735864	0.259845	0.744983	0.800086	0.2465
smoothness_se	-0.222600	0.006614	-0.202694	-0.166777	0.3323
compactness_se	0.206000	0.191975	0.250744	0.212583	0.3189
concavity_se	0.194204	0.143293	0.228082	0.207660	0.2483
concave points_se	0.376169	0.163851	0.407217	0.372320	0.3806
symmetry_se	-0.104321	0.009127	-0.081629	-0.072497	0.2007
fractal_dimension_se	-0.042641	0.054458	-0.005523	-0.019887	0.2836
radius_worst	0.969539	0.352573	0.969476	0.962746	0.2131
texture_worst	0.297008	0.912045	0.303038	0.287489	0.0360
perimeter_worst	0.965137	0.358040	0.970387	0.959120	0.2388
area_worst	0.941082	0.343546	0.941550	0.959213	0.2067
smoothness_worst	0.119616	0.077503	0.150549	0.123523	0.8053
compactness_worst	0.413463	0.277830	0.455774	0.390410	0.4724
concavity_worst	0.526911	0.301025	0.563879	0.512606	0.4349
concave points_worst	0.744214	0.295316	0.771241	0.722017	0.5030
symmetry_worst	0.163953	0.105008	0.189115	0.143570	0.3943
fractal_dimension_worst	0.007066	0.119205	0.051019	0.003738	0.4993

30 rows × 30 columns

We see that none of the columns are highly correlated.

```
In [9]: printmd('## 1.7. Mapping the Categorical Attr.')

# Mapping Benign to 0 and Malignant to 1
data['diagnosis'] = data['diagnosis'].map({'M':1, 'B':0})
#Check the data stats
data.describe()
```

## 1.7. Mapping the Categorical Attr.

```
Out[9]:
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness
<b>count</b>	569.000000	569.000000	569.000000	569.000000	569.000000	569.0
<b>mean</b>	14.127292	19.289649	91.969033	654.889104	0.096360	0.
<b>std</b>	3.524049	4.301036	24.298981	351.914129	0.014064	0.
<b>min</b>	6.981000	9.710000	43.790000	143.500000	0.052630	0.
<b>25%</b>	11.700000	16.170000	75.170000	420.300000	0.086370	0.
<b>50%</b>	13.370000	18.840000	86.240000	551.100000	0.095870	0.
<b>75%</b>	15.780000	21.800000	104.100000	782.700000	0.105300	0.
<b>max</b>	28.110000	39.280000	188.500000	2501.000000	0.163400	0.

8 rows × 30 columns


## 2. Exploratory Analysis

```
In [10]: radius = data[['radius_mean', 'radius_se', 'radius_worst', 'diagnosis']]

# M : 1, B: 0
sns.pairplot(radius, hue='diagnosis', palette="husl", markers=["o", "s"], size=4)
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\axisgrid.py:2076: UserWarning:  
The `size` parameter has been renamed to `height`; please update your code.  
warnings.warn(msg, UserWarning)


```
Out[10]: <seaborn.axisgrid.PairGrid at 0x17d678dc5e0>
```



```
In [11]: sns.pairplot(radius,kind="reg",size=4)
```

```
C:\ProgramData\Anaconda3\lib\site-packages\seaborn\axisgrid.py:2076: UserWarning:  
The `size` parameter has been renamed to `height`; please update your code.  
warnings.warn(msg, UserWarning)
```


```
Out[11]: <seaborn.axisgrid.PairGrid at 0x17d67918610>
```



```
In [12]: g = sns.PairGrid(radius, hue='diagnosis', palette="Set1", size=4)
g = g.map_diag(plt.hist)
g = g.map_offdiag(plt.scatter, s = 3)
```


C:\ProgramData\Anaconda3\lib\site-packages\seaborn\axisgrid.py:1209: UserWarning:  
The `size` parameter has been renamed to `height`; please update your code.  
warnings.warn(UserWarning(msg))





```
In [13]: g = sns.PairGrid(radius, hue="diagnosis", palette="Set2", size=4, hue_kws={"marker":
g = g.map(plt.scatter, linewidths=1, edgecolor="w", s=40)
g = g.add_legend()
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\axisgrid.py:1209: UserWarning:  
The `size` parameter has been renamed to `height`; please update your code.  
warnings.warn(UserWarning(msg))



```
In [14]: printmd('## 2.1. Scaling the Dataset')
from sklearn import preprocessing
# Scaling the dataset
datas = pd.DataFrame(preprocessing.scale(data.iloc[:,1:32]))
datas.columns = list(data.iloc[:,1:32].columns)
datas['diagnosis'] = data['diagnosis']
datas.head()
```

## 2.1. Scaling the Dataset


```
Out[14]:
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean
0	1.097064	-2.073335	1.269934	0.984375	1.568466	3.283511
1	1.829821	-0.353632	1.685955	1.908708	-0.826962	-0.487077
2	1.579888	0.456187	1.566503	1.558884	0.942210	1.052920
3	-0.768909	0.253732	-0.592687	-0.764464	3.283553	3.402909
4	1.750297	-1.151816	1.776573	1.826229	0.280372	0.539340

5 rows × 31 columns

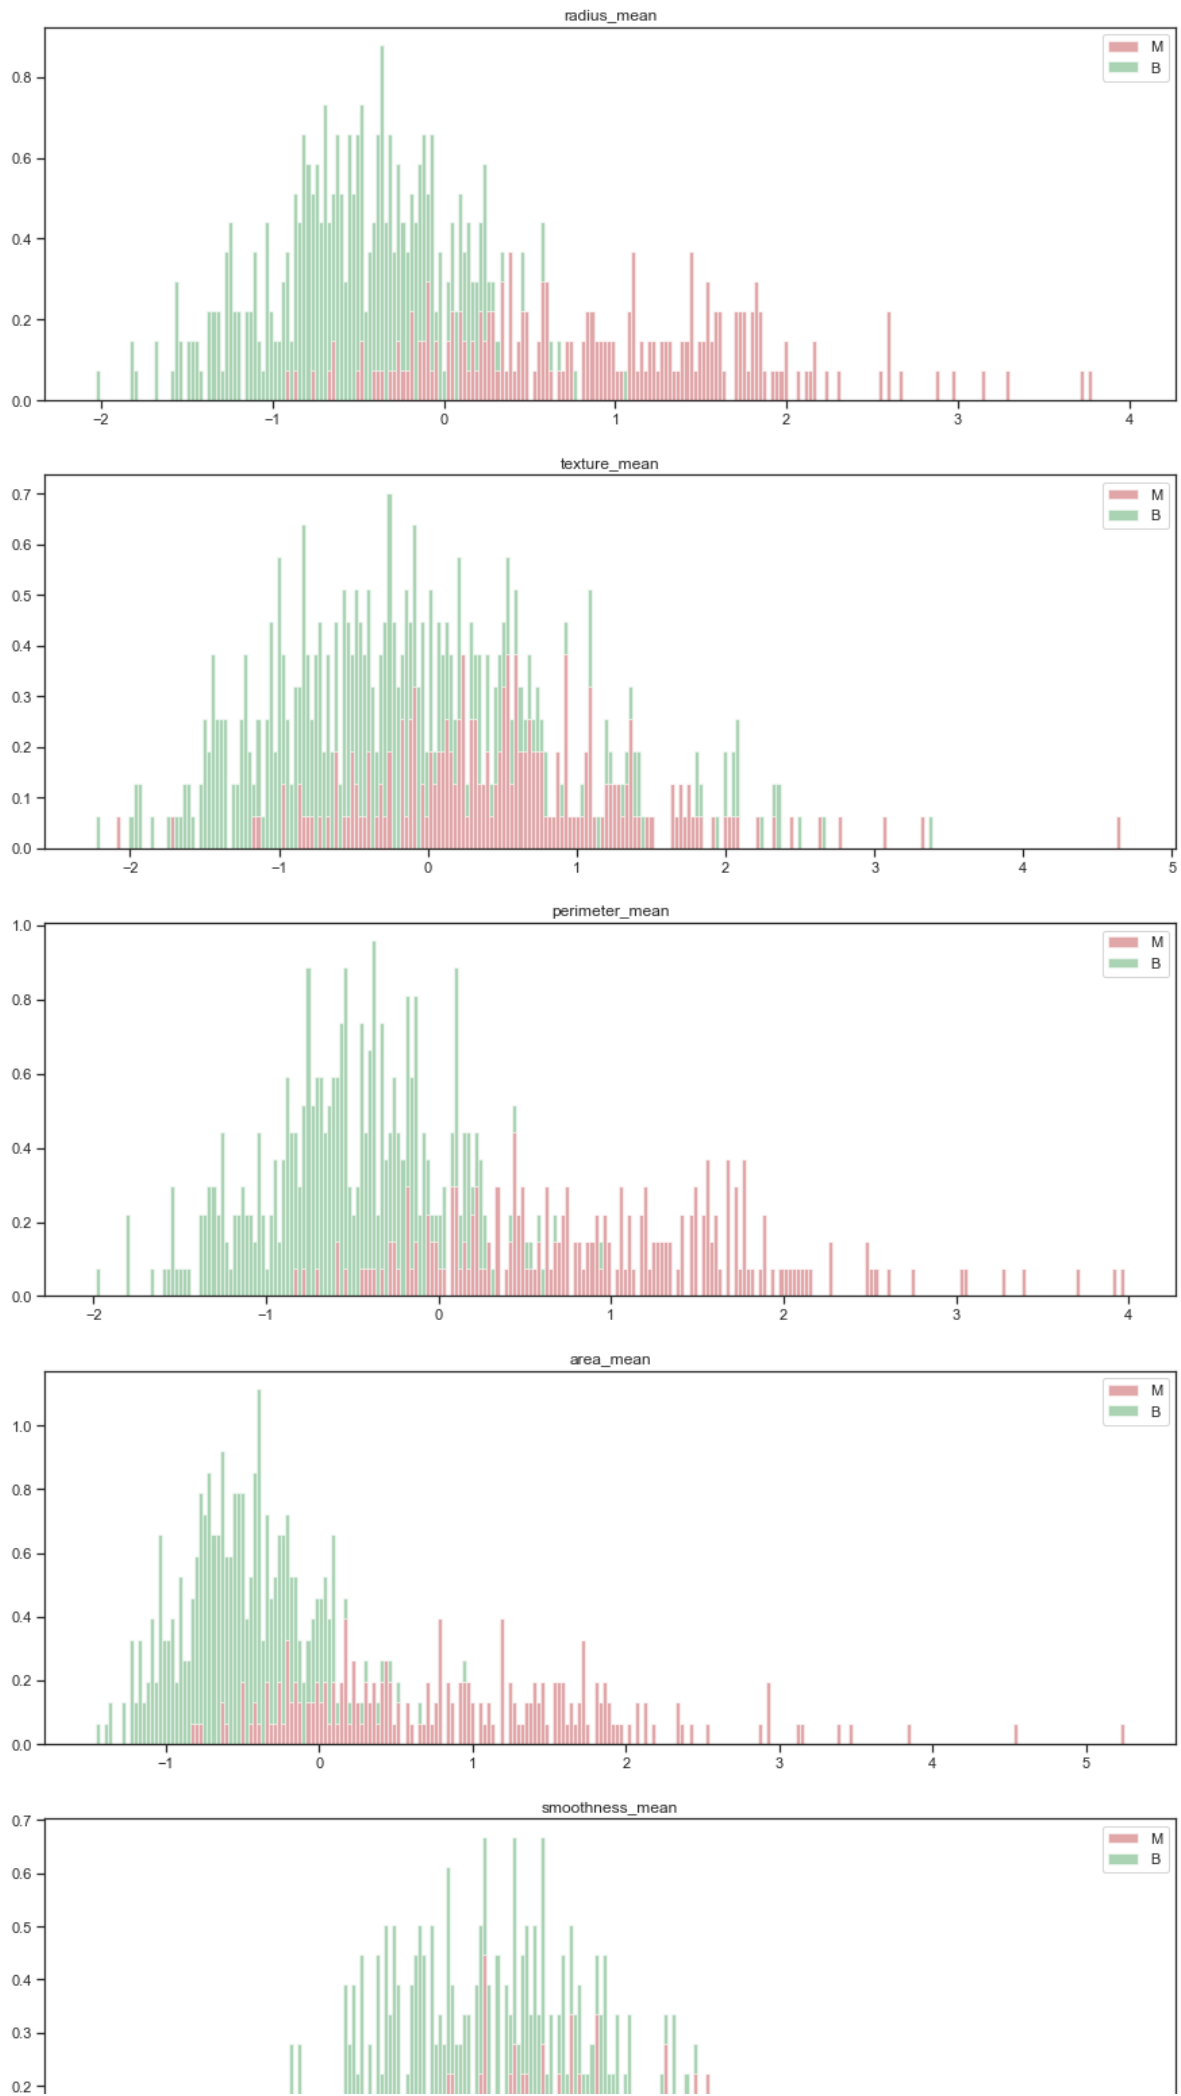
```
In [15]: #draw a heatmap between mean features and diagnosis
features_mean = ['radius_mean', 'texture_mean', 'perimeter_mean', 'area_mean', 'smoothness_mean', 'compactness_mean']
```


```
plt.figure(figsize=(15,15))
heat = sns.heatmap(datas[features_mean].corr(), vmax=1, square=True, annot=True)
```




```
In [16]: import numpy as np
# Splitting the dataset into malignant and benign.
dataMalignant=datas[datas['diagnosis'] ==1]
dataBenign=datas[datas['diagnosis'] ==0]

#Plotting these features as a histogram
fig, axes = plt.subplots(nrows=10, ncols=1, figsize=(15,60))
for idx,ax in enumerate(axes):
    ax.figure
    binwidth= (max(datas[features_mean[idx]]) - min(datas[features_mean[idx]]))/250
    ax.hist([dataMalignant[features_mean[idx]],dataBenign[features_mean[idx]]], binwidth)
    ax.legend(loc='upper right')
    ax.set_title(features_mean[idx])
plt.show()
```








```
In [17]: printmd('## Box plot')
data.select_dtypes(exclude = 'category').plot(kind = 'box', figsize = (10,8))
```

## Box plot

Out[17]: <AxesSubplot:>



## 4. Model Development & Classification

### 4.1. Data Preparation

```
In [18]: data
```

Out[18]:

	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	com
<b>0</b>	1	17.99	10.38	122.80	1001.0	0.11840	
<b>1</b>	1	20.57	17.77	132.90	1326.0	0.08474	
<b>2</b>	1	19.69	21.25	130.00	1203.0	0.10960	
<b>3</b>	1	11.42	20.38	77.58	386.1	0.14250	
<b>4</b>	1	20.29	14.34	135.10	1297.0	0.10030	
...	...	...	...	...	...	...	...
<b>564</b>	1	21.56	22.39	142.00	1479.0	0.11100	
<b>565</b>	1	20.13	28.25	131.20	1261.0	0.09780	
<b>566</b>	1	16.60	28.08	108.30	858.1	0.08455	
<b>567</b>	1	20.60	29.33	140.10	1265.0	0.11780	
<b>568</b>	0	7.76	24.54	47.92	181.0	0.05263	

569 rows × 31 columns

In [19]:

```
X = data.iloc[:, 1:].values
y = data.iloc[:, 0].values

from sklearn.preprocessing import LabelEncoder
le = LabelEncoder()
y = le.fit_transform(y)
```

In [20]:

X

Out[20]:

```
array([[1.799e+01, 1.038e+01, 1.228e+02, ..., 2.654e-01, 4.601e-01,
        1.189e-01],
       [2.057e+01, 1.777e+01, 1.329e+02, ..., 1.860e-01, 2.750e-01,
        8.902e-02],
       [1.969e+01, 2.125e+01, 1.300e+02, ..., 2.430e-01, 3.613e-01,
        8.758e-02],
       ...,
       [1.660e+01, 2.808e+01, 1.083e+02, ..., 1.418e-01, 2.218e-01,
        7.820e-02],
       [2.060e+01, 2.933e+01, 1.401e+02, ..., 2.650e-01, 4.087e-01,
        1.240e-01],
       [7.760e+00, 2.454e+01, 4.792e+01, ..., 0.000e+00, 2.871e-01,
        7.039e-02]])
```

In [21]:

y

```
Out[21]: array([1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 0,
1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, 1, 1, 1,
1, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1,
0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 1, 1,
0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0,
0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1,
1, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 0,
0, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1,
1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 1,
0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0,
0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1,
1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 1, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1,
1, 0, 0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1,
1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 1, 1, 0, 1,
0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0,
0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0,
0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0,
0, 0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0],
dtype=int64)
```

```
In [22]: print(data.info())
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 31 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   diagnosis                             569 non-null    category
1   radius_mean                           569 non-null    float64
2   texture_mean                           569 non-null    float64
3   perimeter_mean                         569 non-null    float64
4   area_mean                             569 non-null    float64
5   smoothness_mean                       569 non-null    float64
6   compactness_mean                      569 non-null    float64
7   concavity_mean                        569 non-null    float64
8   concave points_mean                   569 non-null    float64
9   symmetry_mean                         569 non-null    float64
10  fractal_dimension_mean                 569 non-null    float64
11  radius_se                              569 non-null    float64
12  texture_se                             569 non-null    float64
13  perimeter_se                           569 non-null    float64
14  area_se                               569 non-null    float64
15  smoothness_se                         569 non-null    float64
16  compactness_se                        569 non-null    float64
17  concavity_se                          569 non-null    float64
18  concave points_se                     569 non-null    float64
19  symmetry_se                           569 non-null    float64
20  fractal_dimension_se                   569 non-null    float64
21  radius_worst                          569 non-null    float64
22  texture_worst                         569 non-null    float64
23  perimeter_worst                       569 non-null    float64
24  area_worst                            569 non-null    float64
25  smoothness_worst                      569 non-null    float64
26  compactness_worst                     569 non-null    float64
27  concavity_worst                       569 non-null    float64
28  concave points_worst                   569 non-null    float64
29  symmetry_worst                        569 non-null    float64
30  fractal_dimension_worst                569 non-null    float64
dtypes: category(1), float64(30)
memory usage: 134.2 KB
None
```

```
In [23]: # Splitting the dataset into the Training set and Test set
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.1, random_
```

```
In [24]: #Feature Scaling
from sklearn.preprocessing import StandardScaler
sc = StandardScaler()
X_train = sc.fit_transform(X_train)
X_test = sc.transform(X_test)
```

## 4.2. Model Development

```
In [25]: import keras
from keras.models import Sequential
from keras.layers import Dense, Dropout
```

```
In [26]: # Initialising the ANN
classifier = Sequential()
```

```
In [27]: #first hidden layer
```

```
classifier.add(Dense(units=9, kernel_initializer='he_uniform', activation='relu', input_shape=(9,)))
```

```
In [28]: classifier.add(Dropout(rate=0.1))
```

```
In [29]: #second hidden layer
classifier.add(Dense(units=9, kernel_initializer='he_uniform', activation='relu'))
```

```
In [30]: # last layer or output layer
classifier.add(Dense(units=1, kernel_initializer='glorot_uniform', activation='sigmoid'))
```

```
In [31]: #taking summary of layers
classifier.summary()
```


Model: "sequential"

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 9)	279
dropout (Dropout)	(None, 9)	0
dense_1 (Dense)	(None, 9)	90
dense_2 (Dense)	(None, 1)	10
Total params: 379		
Trainable params: 379		
Non-trainable params: 0		

```
In [32]: from sklearn.decomposition import PCA
pca = PCA(n_components=30)
pca.fit(X_train)

plt.plot(np.cumsum(pca.explained_variance_ratio_))
plt.xlabel('Number of components')
plt.ylabel('Cumulative explained variance')
```

```
Out[32]: Text(0, 0.5, 'Cumulative explained variance')
```



```
In [33]: classifier.compile(optimizer='adam', loss='binary_crossentropy', metrics='accuracy')
```

```
In [34]: X_train
```

```
Out[34]: array([[ -0.52787029,  2.49821982, -0.59939466, ..., -1.74713139,
                -0.79044533, -0.91054389],
                [ -0.55333608,  0.29431013, -0.60759343, ..., -0.62275667,
                -0.33646358, -0.83551633],
                [ 2.15452653,  0.40392257,  2.26525805, ...,  1.03846122,
                -0.11504791,  0.26488788],
                ...,
                [-1.3297598 , -0.21876938, -1.32088704, ..., -0.98271999,
                -0.718764  , -0.13637062],
                [-1.24940108, -0.24209117, -1.2835826 , ..., -1.74713139,
                -1.58690456, -1.01280367],
                [-0.74291476,  1.08958336, -0.71827692, ..., -0.2865488 ,
                -1.26354211,  0.19486216]])
```

```
In [35]: len(X_train)
```

```
Out[35]: 512
```

```
In [36]: y_train
```

```
Out[36]: array([0, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0,
                0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0,
                0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1,
                0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0,
                1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 0, 1, 0, 1, 1,
                1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 1,
                0, 0, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0,
                0, 0, 0, 1, 0, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0,
                1, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 1,
                0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
                0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 0,
                0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0,
                0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1,
                1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0,
                0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0,
                0, 1, 1, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1,
                0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1,
                1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0,
                1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0,
                0, 1, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 1, 0, 0, 0,
                0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0,
                1, 1, 1, 0, 0, 0], dtype=int64)
```

```
In [37]: len(y_train)
```

```
Out[37]: 512
```

```
In [39]: # Fitting the ANN to the Training set
history = classifier.fit(X_train, y_train, batch_size=32, epochs=50, validation_sp
# Long scroll ahead but worth
# The batch size and number of epochs have been set using trial and error. Still lo
```

```
Epoch 1/50
13/13 [=====] - 1s 13ms/step - loss: 0.6886 - accuracy: 0.6381 - val_loss: 0.6246 - val_accuracy: 0.6019
Epoch 2/50
13/13 [=====] - 0s 3ms/step - loss: 0.6014 - accuracy: 0.6699 - val_loss: 0.5362 - val_accuracy: 0.6408
Epoch 3/50
13/13 [=====] - 0s 3ms/step - loss: 0.5387 - accuracy: 0.7066 - val_loss: 0.4654 - val_accuracy: 0.7184
Epoch 4/50
13/13 [=====] - 0s 3ms/step - loss: 0.4715 - accuracy: 0.7677 - val_loss: 0.4013 - val_accuracy: 0.7961
Epoch 5/50
13/13 [=====] - 0s 3ms/step - loss: 0.4360 - accuracy: 0.8117 - val_loss: 0.3484 - val_accuracy: 0.8932
Epoch 6/50
13/13 [=====] - 0s 3ms/step - loss: 0.3941 - accuracy: 0.8337 - val_loss: 0.3057 - val_accuracy: 0.9126
Epoch 7/50
13/13 [=====] - 0s 3ms/step - loss: 0.3653 - accuracy: 0.8460 - val_loss: 0.2712 - val_accuracy: 0.9223
Epoch 8/50
13/13 [=====] - 0s 3ms/step - loss: 0.3191 - accuracy: 0.8729 - val_loss: 0.2433 - val_accuracy: 0.9709
Epoch 9/50
13/13 [=====] - 0s 3ms/step - loss: 0.2982 - accuracy: 0.8924 - val_loss: 0.2178 - val_accuracy: 0.9806
Epoch 10/50
13/13 [=====] - 0s 3ms/step - loss: 0.2650 - accuracy: 0.9169 - val_loss: 0.1950 - val_accuracy: 0.9806
Epoch 11/50
13/13 [=====] - 0s 3ms/step - loss: 0.2429 - accuracy: 0.9315 - val_loss: 0.1735 - val_accuracy: 0.9903
Epoch 12/50
13/13 [=====] - 0s 3ms/step - loss: 0.2273 - accuracy: 0.9413 - val_loss: 0.1553 - val_accuracy: 0.9903
Epoch 13/50
13/13 [=====] - 0s 3ms/step - loss: 0.2134 - accuracy: 0.9413 - val_loss: 0.1401 - val_accuracy: 0.9903
Epoch 14/50
13/13 [=====] - 0s 3ms/step - loss: 0.1927 - accuracy: 0.9462 - val_loss: 0.1281 - val_accuracy: 0.9903
Epoch 15/50
13/13 [=====] - 0s 3ms/step - loss: 0.1796 - accuracy: 0.9560 - val_loss: 0.1181 - val_accuracy: 0.9903
Epoch 16/50
13/13 [=====] - 0s 3ms/step - loss: 0.1689 - accuracy: 0.9511 - val_loss: 0.1092 - val_accuracy: 0.9903
Epoch 17/50
13/13 [=====] - 0s 3ms/step - loss: 0.1534 - accuracy: 0.9584 - val_loss: 0.1013 - val_accuracy: 0.9903
Epoch 18/50
13/13 [=====] - 0s 3ms/step - loss: 0.1547 - accuracy: 0.9560 - val_loss: 0.0955 - val_accuracy: 0.9903
Epoch 19/50
13/13 [=====] - 0s 3ms/step - loss: 0.1536 - accuracy: 0.9511 - val_loss: 0.0909 - val_accuracy: 0.9903
Epoch 20/50
13/13 [=====] - 0s 3ms/step - loss: 0.1322 - accuracy: 0.9584 - val_loss: 0.0861 - val_accuracy: 0.9903
Epoch 21/50
13/13 [=====] - 0s 3ms/step - loss: 0.1269 - accuracy: 0.9584 - val_loss: 0.0817 - val_accuracy: 0.9903
Epoch 22/50
```

```
13/13 [=====] - 0s 3ms/step - loss: 0.1289 - accuracy: 0.9584 - val_loss: 0.0781 - val_accuracy: 0.9903
Epoch 23/50
13/13 [=====] - 0s 3ms/step - loss: 0.1154 - accuracy: 0.9658 - val_loss: 0.0747 - val_accuracy: 0.9903
Epoch 24/50
13/13 [=====] - 0s 3ms/step - loss: 0.1163 - accuracy: 0.9658 - val_loss: 0.0720 - val_accuracy: 0.9903
Epoch 25/50
13/13 [=====] - 0s 3ms/step - loss: 0.1081 - accuracy: 0.9707 - val_loss: 0.0692 - val_accuracy: 0.9903
Epoch 26/50
13/13 [=====] - 0s 3ms/step - loss: 0.1137 - accuracy: 0.9658 - val_loss: 0.0672 - val_accuracy: 0.9903
Epoch 27/50
13/13 [=====] - 0s 3ms/step - loss: 0.1096 - accuracy: 0.9658 - val_loss: 0.0645 - val_accuracy: 0.9903
Epoch 28/50
13/13 [=====] - 0s 3ms/step - loss: 0.1037 - accuracy: 0.9658 - val_loss: 0.0623 - val_accuracy: 0.9903
Epoch 29/50
13/13 [=====] - 0s 3ms/step - loss: 0.0974 - accuracy: 0.9731 - val_loss: 0.0602 - val_accuracy: 0.9903
Epoch 30/50
13/13 [=====] - 0s 3ms/step - loss: 0.1022 - accuracy: 0.9731 - val_loss: 0.0586 - val_accuracy: 0.9903
Epoch 31/50
13/13 [=====] - 0s 3ms/step - loss: 0.0977 - accuracy: 0.9731 - val_loss: 0.0567 - val_accuracy: 0.9903
Epoch 32/50
13/13 [=====] - 0s 3ms/step - loss: 0.0922 - accuracy: 0.9731 - val_loss: 0.0546 - val_accuracy: 0.9903
Epoch 33/50
13/13 [=====] - 0s 3ms/step - loss: 0.0880 - accuracy: 0.9756 - val_loss: 0.0528 - val_accuracy: 0.9903
Epoch 34/50
13/13 [=====] - 0s 3ms/step - loss: 0.0901 - accuracy: 0.9756 - val_loss: 0.0514 - val_accuracy: 0.9903
Epoch 35/50
13/13 [=====] - 0s 3ms/step - loss: 0.0872 - accuracy: 0.9756 - val_loss: 0.0502 - val_accuracy: 0.9903
Epoch 36/50
13/13 [=====] - 0s 3ms/step - loss: 0.0844 - accuracy: 0.9780 - val_loss: 0.0487 - val_accuracy: 0.9903
Epoch 37/50
13/13 [=====] - 0s 3ms/step - loss: 0.0832 - accuracy: 0.9756 - val_loss: 0.0465 - val_accuracy: 0.9903
Epoch 38/50
13/13 [=====] - 0s 3ms/step - loss: 0.0913 - accuracy: 0.9731 - val_loss: 0.0454 - val_accuracy: 0.9903
Epoch 39/50
13/13 [=====] - 0s 3ms/step - loss: 0.0896 - accuracy: 0.9780 - val_loss: 0.0443 - val_accuracy: 0.9903
Epoch 40/50
13/13 [=====] - 0s 3ms/step - loss: 0.0856 - accuracy: 0.9731 - val_loss: 0.0431 - val_accuracy: 0.9903
Epoch 41/50
13/13 [=====] - 0s 3ms/step - loss: 0.0850 - accuracy: 0.9780 - val_loss: 0.0419 - val_accuracy: 0.9903
Epoch 42/50
13/13 [=====] - 0s 3ms/step - loss: 0.0827 - accuracy: 0.9756 - val_loss: 0.0410 - val_accuracy: 0.9903
Epoch 43/50
13/13 [=====] - 0s 3ms/step - loss: 0.0752 - accuracy: 0.
```

```

9804 - val_loss: 0.0397 - val_accuracy: 0.9903
Epoch 44/50
13/13 [=====] - 0s 3ms/step - loss: 0.0843 - accuracy: 0.
9804 - val_loss: 0.0389 - val_accuracy: 0.9903
Epoch 45/50
13/13 [=====] - 0s 3ms/step - loss: 0.0772 - accuracy: 0.
9780 - val_loss: 0.0375 - val_accuracy: 0.9903
Epoch 46/50
13/13 [=====] - 0s 3ms/step - loss: 0.0761 - accuracy: 0.
9780 - val_loss: 0.0364 - val_accuracy: 1.0000
Epoch 47/50
13/13 [=====] - 0s 3ms/step - loss: 0.0766 - accuracy: 0.
9756 - val_loss: 0.0350 - val_accuracy: 1.0000
Epoch 48/50
13/13 [=====] - 0s 5ms/step - loss: 0.0711 - accuracy: 0.
9829 - val_loss: 0.0342 - val_accuracy: 1.0000
Epoch 49/50
13/13 [=====] - 0s 3ms/step - loss: 0.0671 - accuracy: 0.
9780 - val_loss: 0.0331 - val_accuracy: 1.0000
Epoch 50/50
13/13 [=====] - 0s 3ms/step - loss: 0.0715 - accuracy: 0.
9804 - val_loss: 0.0322 - val_accuracy: 1.0000

```

```

In [40]: # Predicting the Test set results
y_pred = classifier.predict(X_test)
y_pred = (y_pred > 0.5)

```

```

In [41]: # Making the Confusion Matrix
from sklearn.metrics import confusion_matrix
cm = confusion_matrix(y_test, y_pred)

```

```

In [42]: print("Our accuracy is {}".format(((cm[0][0] + cm[1][1])/57)*100))


Our accuracy is 96.49122807017544%

```

```

In [43]: sns.heatmap(cm,annot=True)
plt.savefig('h.png')

```



```

In [ ]:


```

```


In [44]: # summarize history for accuracy
plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('model accuracy')
plt.ylabel('accuracy')
plt.xlabel('epoch')

```

```
plt.legend(['train', 'test'], loc='lower right')
plt.show()
```



```
In [45]: # summarize history for loss
plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('model loss')
plt.ylabel('loss')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='lower left')
plt.show()
```



```
In [46]: classifier.evaluate(X_test, y_test, verbose=2)
```

```
2/2 - 0s - loss: 0.0750 - accuracy: 0.9649 - 14ms/epoch - 7ms/step
[0.07502038776874542, 0.9649122953414917]
```

```
Out[46]:
```

```
In [47]: y_pred = classifier.predict(X_test, verbose=0)
y_pred
```

```
Out[47]: array([[9.8528194e-01],
               [4.2927146e-02],
               [6.1172247e-04],
               [6.7895651e-03],
               [8.1724393e-05],
               [2.8225183e-03],
               [7.2002498e-05],
               [5.4946542e-04],
               [1.9304873e-05],
               [7.8006798e-07],
               [3.0570233e-01],
               [1.0753912e-01],
               [7.3306014e-06],
               [6.5287519e-01],
               [8.2655454e-01],
               [9.9573505e-01],
               [2.2798777e-04],
               [9.9981570e-01],
               [9.9914670e-01],
               [9.9999940e-01],
               [9.9596143e-01],
               [9.9139106e-01],
               [4.4696629e-03],
               [7.5700879e-04],
               [9.9949783e-01],
               [1.9466877e-04],
               [8.5376541e-06],
               [9.9427080e-01],
               [7.2956085e-04],
               [9.9994797e-01],
               [1.0338121e-05],
               [9.9951959e-01],
               [9.8506540e-02],
               [9.7987396e-01],
               [3.8287490e-07],
               [9.6691626e-01],
               [9.4931126e-03],
               [9.9814558e-01],
               [4.9524903e-03],
               [9.8646605e-01],
               [7.5975442e-01],
               [1.3745483e-05],
               [7.3339814e-01],
               [2.0404559e-05],
               [5.3348005e-02],
               [9.9999940e-01],
               [1.4404962e-07],
               [3.0549169e-02],
               [4.0635467e-04],
               [9.9941981e-01],
               [9.9986076e-01],
               [9.7790194e-01],
               [9.9879289e-01],
               [6.6787004e-04],
               [1.5106201e-03],
               [1.1739471e-04],
               [1.3407469e-03]], dtype=float32)
```

```
In [49]: from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred.round(), labels=[0,1]))
```



	precision	recall	f1-score	support
0	1.00	0.94	0.97	35
1	0.92	1.00	0.96	22
accuracy			0.96	57
macro avg	0.96	0.97	0.96	57
weighted avg	0.97	0.96	0.97	57

In [ ]: